



Performance Analysis of Cilium

An eBPF based CNI

- 1) Anvaya Bheemanakone Narappa
- 2) Mohammed Syed Akbar Hashmi

Under supervision of Federico Parola



Problem Statement

To perform the performance of analysis of Cilium (ebpf based CNI) and Calico (standard networking - kube-proxy) with respect to number of Layer 3/4 Layer 7 policies.

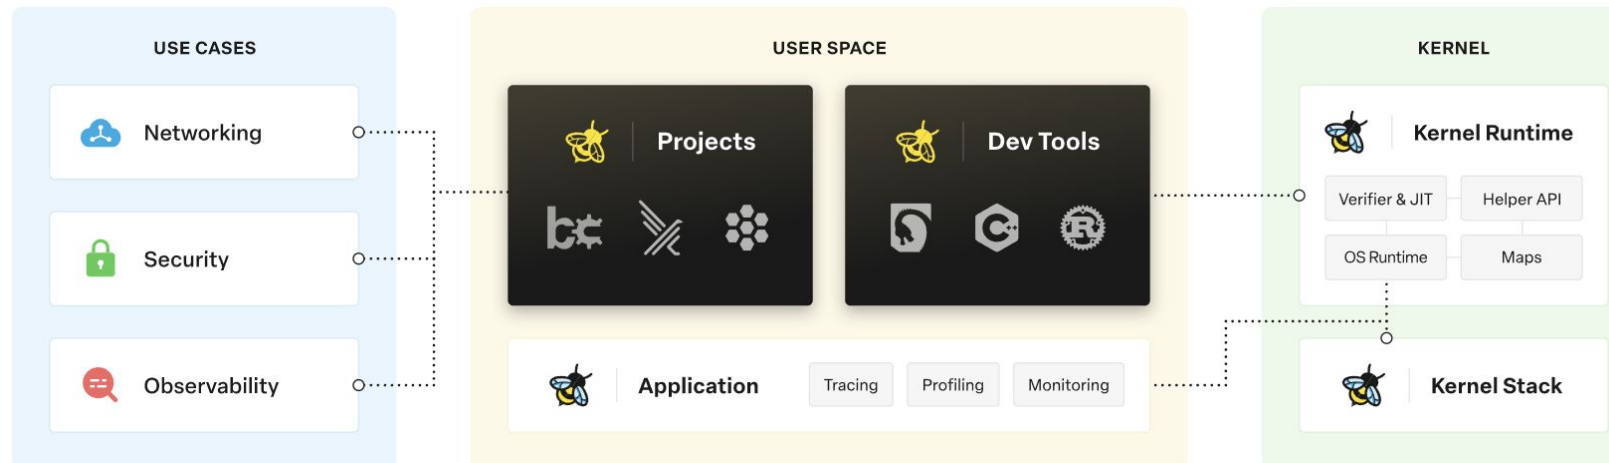
eBPF

eBPF (extended Berkeley Packet Filter) is a powerful and flexible technology that allows users to run custom programs within the Linux kernel. It was initially developed for network packet filtering and analysis, but has since been extended to other areas such as tracing and performance analysis.

Why eBPF?

- **Performance:** eBPF drastically improves processing by being JIT compiled and running directly in the kernel.
- **Security:** eBPF programs are verified to not crash the kernel and can only be modified by privileged users.
- **Flexibility:** Modify or add functionality and use cases to the kernel without having to restart or patch it.

- eBPF programs can be hooked anywhere on the kernel to modify functionality.
- Programs are verified to execute safely on the kernel
- Many applications such as bcc, Cilium, Falco, Katran use ebpf.
- The Use cases include Networking, Security and Observability.

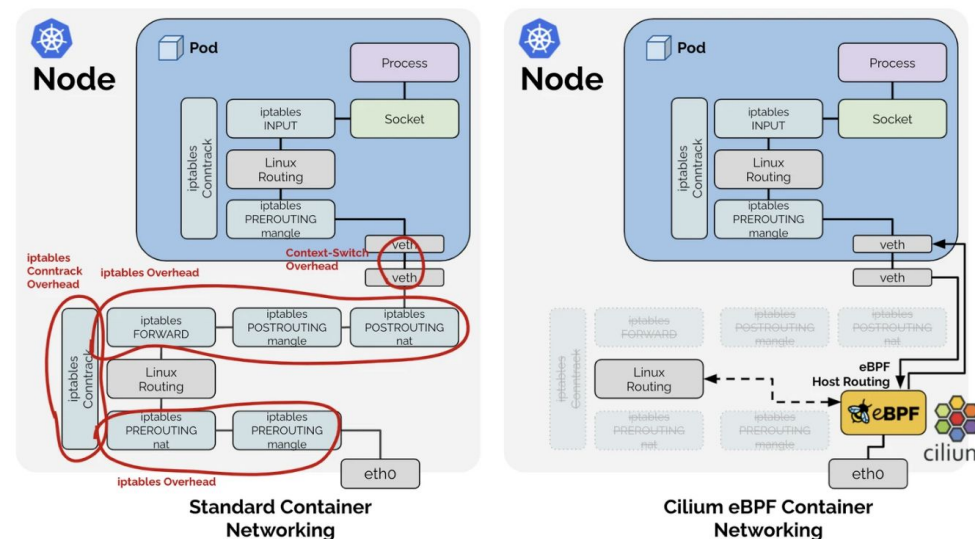


Cilium

- Open-source software-defined networking (SDN) solution
- Provides secure and efficient networking for containerized applications running in Kubernetes environments.
- Built on top of Linux kernel technologies such as eBPF (extended Berkeley Packet Filter) and XDP (eXpress Data Path)
- Cilium offers several features that help improve networking in Kubernetes environments such as High performance, Secure connectivity, Observability

Container Networking (Standard vs Cilium ebpf)

- ebpf host routing in Cilium Allows to bypass all the iptables overhead and as well as some context switching overhead while traversing through the veth pair.
- Network packets are picked up as early as possible from the network device facing the network and delivered directly into the network namespace of the Kubernetes Pod



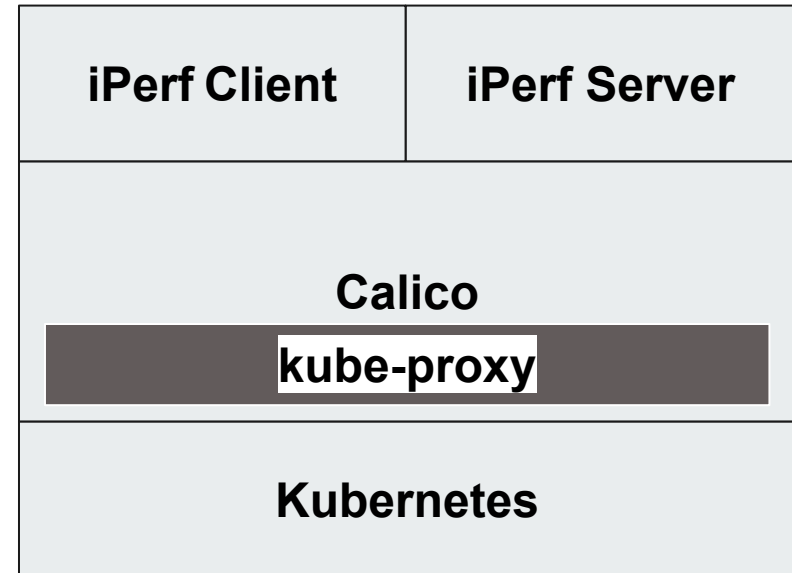
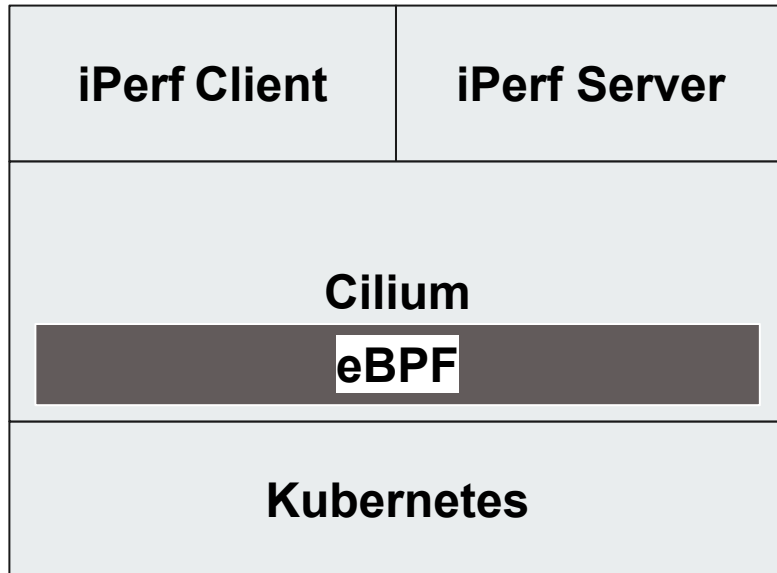
Tools/Frameworks

- Kubernetes
- GCP
- Cilium CNI
- Calico
- iperf/netperf
- Apache benchmark

Environment Setup

- Hosted on Google Cloud Platform
- Linux v5.4.0-1101-gcp -> Ubuntu 18.04.1
- 2 CPU Cores - Intel(R) Xeon(R) CPU @ 2.20GHz
- Cilium v1.13.0
- Calico v3.24.5
- NIC - 10Gbps
- Kubernetes v1.19 (1 Master +1 Worker)

Experiment Setup



Parameter Tuning for Cilium

We tuned the environment to achieve maximum benchmark

The parameters that were changed were

- Disabled Hubble as this added an additional overhead of 1-15 %
- Increased the MTU size - the system was configured to use jumbo size frames
- Kube-Proxy replacement was set to strict

```
an001@workercilium1:~$ ifconfig
cilium_host: flags=4291<UP,BROADCAST,RUNNING,NOARP,MULTICAST> mtu 65000
    inet 10.0.1.218 netmask 255.255.255.255 broadcast 0.0.0.0
    inet6 fe80::c40f:ecff:fe63:af0d prefixlen 64 scopeid 0x20<link>
    ether c6:0f:ec:63:af:0d txqueuelen 1000 (Ethernet)
    RX packets 38446 bytes 2965992 (2.9 MB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 124 bytes 8752 (8.7 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

cilium_net: flags=4291<UP,BROADCAST,RUNNING,NOARP,MULTICAST> mtu 65000
    inet6 fe80::d02d:92ff:feb4:6e40 prefixlen 64 scopeid 0x20<link>
    ether d2:2d:92:b4:6e:40 txqueuelen 1000 (Ethernet)
    RX packets 124 bytes 8752 (8.7 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 38446 bytes 2965992 (2.9 MB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

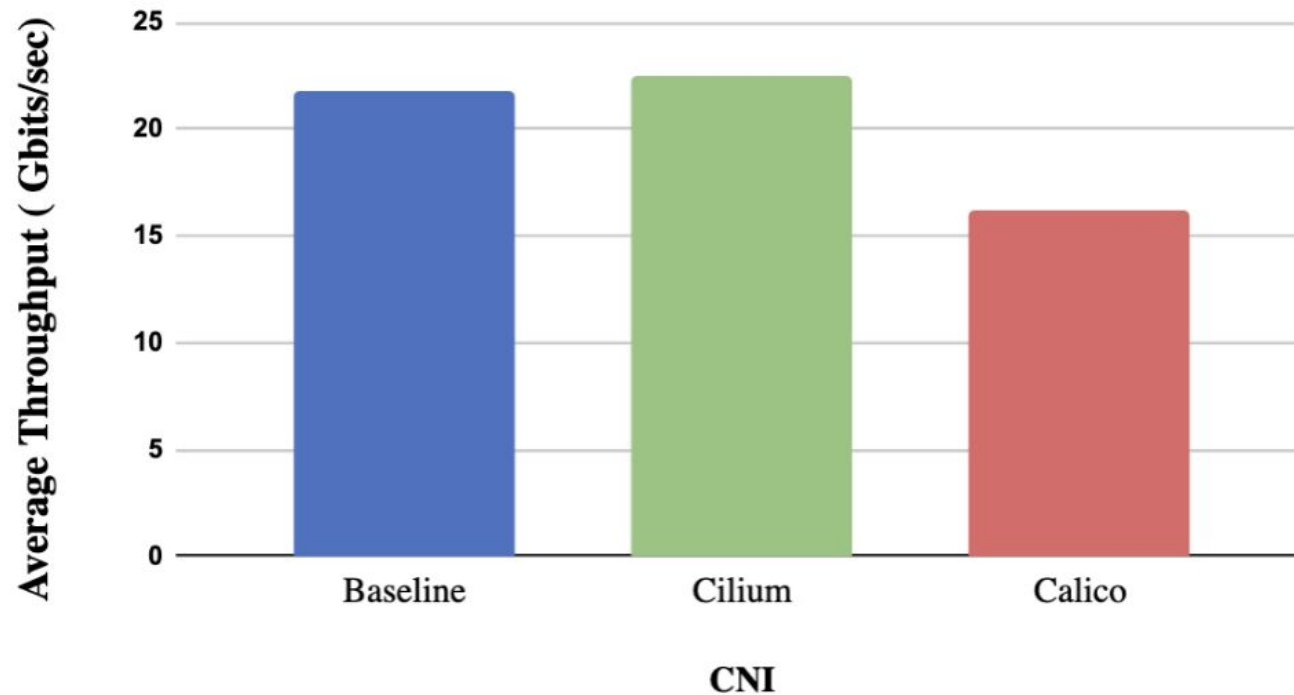
```
an001@mastercilium:~/CiliumProject/apachebenchmark$ helm install cilium cilium/cilium --version 1.13.0 \
> --namespace kube-system \
> --set kubeProxyReplacement=strict \
> --set hubble.enabled=false
```

Experiments Performed

1. General benchmarks for TCP/UDP - Provides baseline
 - a. Throughput - refers to the amount of data that can be transmitted over a network connection within a certain time period
2. Layer3/4 policy testing for TCP/UDP
 - a. cpu utilisation - refers to the percentage of time that the CPU is busy processing instructions. It is used to evaluate the efficiency and capacity of a network connection or network device
3. L7 Policy Testing
 - a. Time per request - measures the time taken to process a single request at the application layer

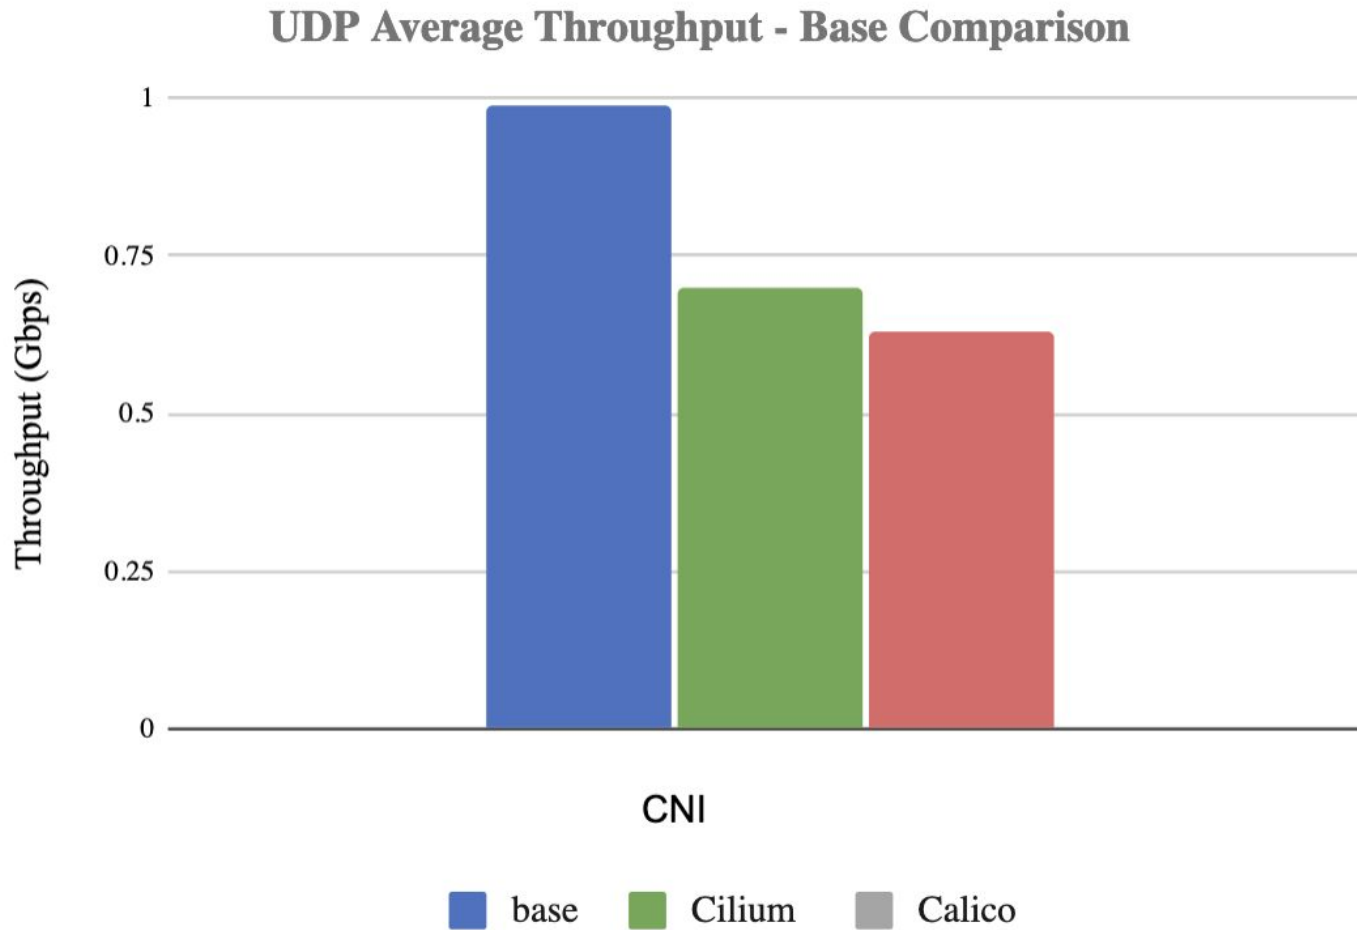
General Benchmarks - TCP

TCP Average Throughput - Base Comparison



- Baseline model doesn't have any CNI for routing
- No networking policies were applied

General Benchmarks - UDP



- tested with bandwidth tuned to 1 Gbps
- Cilium still outperforms Calico

Layer3/4 policy testing for TCP

```
CiliumProject > CiliumNetworkPolicy > $ CiliumpolicygeneratorLabel.sh
1  #!/bin/bash
2  v="v300"
3  x=0
4  NUMBER=$1
5  rm ciliumbasepolicyv222.yaml
6  kubectl delete cnf iperfpolicy$v
7  echo "kind: CiliumNetworkPolicy" >> ciliumbasepolicyv222.yaml
8  echo "apiVersion: \"cilium.io/v2\"" >> ciliumbasepolicyv222.yaml
9  echo "metadata:" >> ciliumbasepolicyv222.yaml
10 echo "  name: \"iperfpolicy$v\"" >> ciliumbasepolicyv222.yaml
11 echo "  namespace: default" >> ciliumbasepolicyv222.yaml
12 echo "spec:" >> ciliumbasepolicyv222.yaml
13 echo "  endpointSelector:" >> ciliumbasepolicyv222.yaml
14 echo "    matchLabels:" >> ciliumbasepolicyv222.yaml
15 echo "      app: iperf" >> ciliumbasepolicyv222.yaml
16 echo "      role: client" >> ciliumbasepolicyv222.yaml
17 echo "      role: server" >> ciliumbasepolicyv222.yaml
18 echo "  ingress:" >> ciliumbasepolicyv222.yaml
19 echo "    - fromEndpoints:" >> ciliumbasepolicyv222.yaml
20 echo "      - matchLabels:" >> ciliumbasepolicyv222.yaml
21 echo "        app: iperf" >> ciliumbasepolicyv222.yaml
22 echo "        role: client" >> ciliumbasepolicyv222.yaml
23
24 while [ $x -le $NUMBER ]; do
25
26   echo "    - matchLabels:" >> ciliumbasepolicyv222.yaml
27   echo "      key$x: value$x" >> ciliumbasepolicyv222.yaml
28   echo "      ((x++))"
29   echo "      key$x: value$x" >> ciliumbasepolicyv222.yaml
30   echo "      ((x++))"
31 done
32
33 ##### Uploading Policy
34
35 echo "Uploading Policy !!!!!!!!!!!!!!!"
36
37 kubectl create -f ciliumbasepolicyv222.yaml
38
39 wait
40 sleep 3
41
42 exit
43
44
45
```

```
CiliumProject > CiliumNetworkPolicy > ! ciliumbasepolicyv222.yaml
1  kind: CiliumNetworkPolicy
2  apiVersion: "cilium.io/v2"
3  metadata:
4    name: "iperfpolicyv300"
5    namespace: default
6  spec:
7    endpointSelector:
8      matchLabels:
9        app: iperf
10       role: client
11       role: server
12    ingress:
13      - fromEndpoints:
14        - matchLabels:
15          app: iperf
16          role: client
17        - matchLabels:
18          key0: value0
19          key1: value1
20        - matchLabels:
21          key2: value2
22          key3: value3
23        - matchLabels:
24          key4: value4
25          key5: value5
26        - matchLabels:
27          key6: value6
28          key7: value7
29        - matchLabels:
30          key8: value8
31          key9: value9
32        - matchLabels:
33          key10: value10
34          key11: value11
35        - matchLabels:
36          key12: value12
37          key13: value13
38        - matchLabels:
39          key14: value14
40          key15: value15
41        - matchLabels:
42          key16: value16
43          key17: value17
44        - matchLabels:
45          key18: value18
46          key19: value19
47        - matchLabels:
48          key20: value20
49          key21: value21
50        - matchLabels:
51          key22: value22
```



```

12 #echo "MTU_SIZE is $MTU_SIZE"
13 echo "Creating Iperf3 pods for stats"
14 kubectl apply -f iperf3.yaml
15 sleep 5
16 echo "Iperf3 pods are created"
17 echo "kubectl get all | grep iperf3 "
18
19 PODNAME=$(kubectl get pods -o wide | gre
20 echo "Podname --> $PODNAME"
21
22 IPERF_SERVER=$(kubectl get pods -o wide
23 echo "iperfServer --> $IPERF_SERVER "
24
25 HOSTIP=$(kubectl get pod $IPERF_SERVER -
26 echo "HostIP --> $HOSTIP"
27 echo "$DATE"
28
29 ##### TCP TEST
30 for i in {1..5}
31 do
32 kubectl exec -it $PODNAME -- iperf3 -c $
33 sleep 2
34 done
35
36
37 # ##### UDP TEST
38 for i in {1..6}
39 do
40 kubectl exec -it $PODNAME -- iperf3 -c $
41 done
42
43 ##### Single UDP test ###
44 # kubectl exec -it $PODNAME -- iperf3 -c
45
46 echo "#####
47 echo "#####
48 echo "Completed the Iperf Test"
49 echo "#####
50 echo "#####
51 # kubectl delete deployment.apps/iperf-c
52 sleep 5
53 # echo "Deleting the iperf3 container"
54 exit
55

```

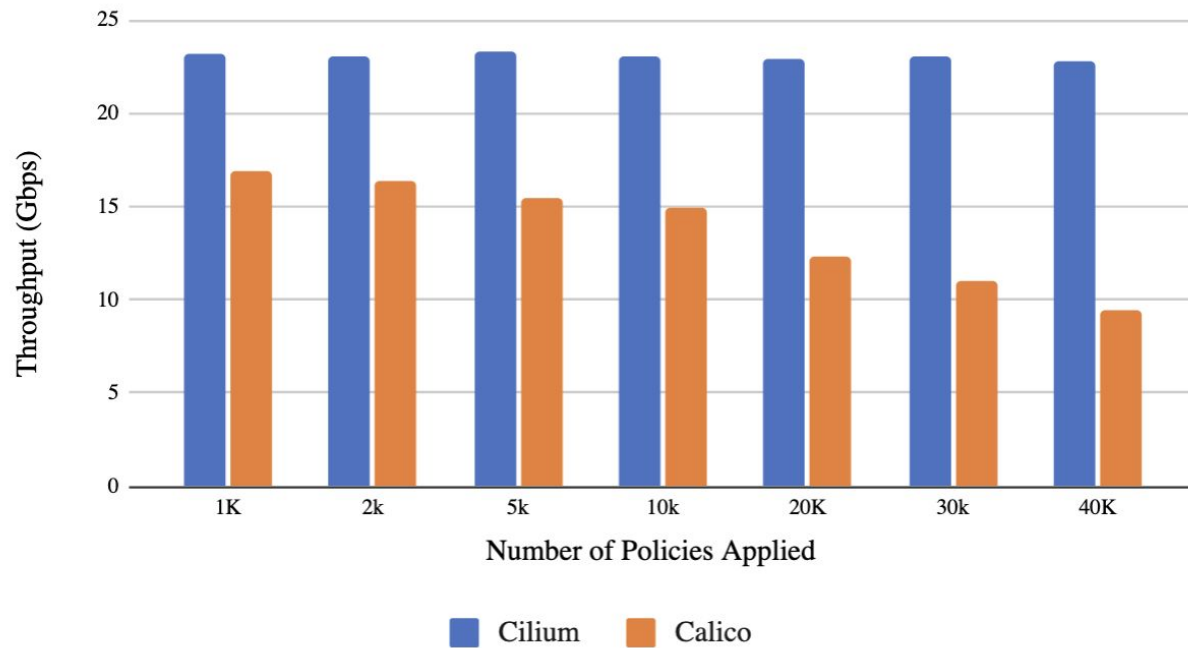
```

iperf 3.9
calicopodtestTCP: calicopodtestTCP: Linux iperf-client-7cc8c447c5-6cv6d 5.4.0-1100-gcp #109~18.04.1-Ubuntu SMP Wed Jan 25 21:16:55 UTC 2023 x86_64
Control connection MSS 1388
calicopodtestTCP: Time: Wed, 22 Mar 2023 03:33:16 GMT
calicopodtestTCP: Connecting to host 192.168.190.254, port 5201
calicopodtestTCP: Cookie: kqd5lu6sjttwflxb22g5rLx5cb2dc6afscrb
calicopodtestTCP: TCP MSS: 1388 (default)
calicopodtestTCP: [ 5] local 192.168.190.255 port 51084 connected to 192.168.190.254 port 5201
calicopodtestTCP: Starting Test: protocol: TCP, 1 streams, 131072 byte blocks, omitting 0 seconds, 10 second test, tos 0
calicopodtestTCP: [ ID] Interval      Transfer      Bitrate      Retr  Cwnd
calicopodtestTCP: [ 5] 0.00-1.00   sec  1.93 GBytes  16.6 Gbits/sec  18  1.15 MBytes
calicopodtestTCP: [ 5] 1.00-2.00   sec   828 MBytes  6.94 Gbits/sec   0  1.18 MBytes
calicopodtestTCP: [ 5] 2.00-3.00   sec  1.72 GBytes  14.8 Gbits/sec   0  1.20 MBytes
calicopodtestTCP: [ 5] 3.00-4.00   sec  2.01 GBytes  17.3 Gbits/sec   0  1.22 MBytes
calicopodtestTCP: [ 5] 4.00-5.00   sec  2.00 GBytes  17.2 Gbits/sec   0  1.24 MBytes
calicopodtestTCP: [ 5] 5.00-6.00   sec  2.03 GBytes  17.4 Gbits/sec   0  1.26 MBytes
calicopodtestTCP: [ 5] 6.00-7.00   sec  1.82 GBytes  15.6 Gbits/sec   0  1.31 MBytes
calicopodtestTCP: [ 5] 7.00-8.00   sec  1.88 GBytes  16.1 Gbits/sec   0  1.34 MBytes
calicopodtestTCP: [ 5] 8.00-9.00   sec  1.89 GBytes  16.2 Gbits/sec   0  1.38 MBytes
calicopodtestTCP: [ 5] 9.00-10.00  sec  2.03 GBytes  17.5 Gbits/sec  392  1.05 MBytes
calicopodtestTCP: -----
calicopodtestTCP: Test Complete. Summary Results:
calicopodtestTCP: [ ID] Interval      Transfer      Bitrate      Retr
calicopodtestTCP: [ 5] 0.00-10.00  sec  18.1 GBytes  15.6 Gbits/sec  410
calicopodtestTCP: [ 5] 0.00-10.00  sec  18.1 GBytes  15.6 Gbits/sec
calicopodtestTCP: CPU Utilization: local/sender 87.8% (2.0%u/85.8%u), remote/receiver 0.1% (0.0%u/0.1%u)
calicopodtestTCP: snd_tcp_congestion cubic
calicopodtestTCP: rcv_tcp_congestion cubic
calicopodtestTCP: iperf Done.
calicopodtestTCP: iperf 3.9
calicopodtestTCP: calicopodtestTCP: Linux iperf-client-7cc8c447c5-6cv6d 5.4.0-1100-gcp #109~18.04.1-Ubuntu SMP Wed Jan 25 21:16:55 UTC 2023 x86_64
Control connection MSS 1388
calicopodtestTCP: Time: Wed, 22 Mar 2023 03:33:28 GMT
calicopodtestTCP: Connecting to host 192.168.190.254, port 5201
calicopodtestTCP: Cookie: oxt3k7cndvfmavf2sxsiu7ssjuwq3fj4h37
calicopodtestTCP: TCP MSS: 1388 (default)
calicopodtestTCP: [ 5] local 192.168.190.255 port 60216 connected to 192.168.190.254 port 5201
calicopodtestTCP: Starting Test: protocol: TCP, 1 streams, 131072 byte blocks, omitting 0 seconds, 10 second test, tos 0
calicopodtestTCP: [ ID] Interval      Transfer      Bitrate      Retr  Cwnd
calicopodtestTCP: [ 5] 0.00-1.01   sec  1.37 GBytes  11.7 Gbits/sec   1  1.39 MBytes
calicopodtestTCP: [ 5] 1.01-2.01   sec  2.09 GBytes  17.9 Gbits/sec   0  1.56 MBytes
calicopodtestTCP: [ 5] 2.01-3.01   sec  2.09 GBytes  18.0 Gbits/sec  453  809 KBytes
calicopodtestTCP: [ 5] 3.01-4.01   sec  1.90 GBytes  16.4 Gbits/sec  11  699 KBytes
calicopodtestTCP: [ 5] 4.01-5.01   sec  1.86 GBytes  16.0 Gbits/sec   0  758 KBytes
calicopodtestTCP: [ 5] 5.01-6.01   sec  1.85 GBytes  15.9 Gbits/sec   0  863 KBytes
calicopodtestTCP: [ 5] 6.01-7.01   sec  2.02 GBytes  17.3 Gbits/sec   0  897 KBytes
calicopodtestTCP: [ 5] 7.01-8.01   sec  2.06 GBytes  17.7 Gbits/sec  92  712 KBytes
calicopodtestTCP: [ 5] 8.01-9.01   sec  2.04 GBytes  17.5 Gbits/sec   0  785 KBytes
calicopodtestTCP: [ 5] 9.01-10.07  sec  1.25 GBytes  10.1 Gbits/sec   0  830 KBytes
calicopodtestTCP: -----
calicopodtestTCP: Test Complete. Summary Results:
calicopodtestTCP: [ ID] Interval      Transfer      Bitrate      Retr
calicopodtestTCP: [ 5] 0.00-10.07  sec  18.5 GBytes  15.8 Gbits/sec  557
calicopodtestTCP: [ 5] 0.00-10.12  sec  18.5 GBytes  15.7 Gbits/sec
calicopodtestTCP: CPU Utilization: local/sender 86.5% (2.1%u/84.4%u), remote/receiver 63.7% (4.5%u/59.2%u)
calicopodtestTCP: snd_tcp_congestion cubic
calicopodtestTCP: rcv_tcp_congestion cubic
calicopodtestTCP: iperf Done.
calicopodtestTCP: iperf 3.9
calicopodtestTCP: calicopodtestTCP: Linux iperf-client-7cc8c447c5-6cv6d 5.4.0-1100-gcp #109~18.04.1-Ubuntu SMP Wed Jan 25 21:16:55 UTC 2023 x86_64

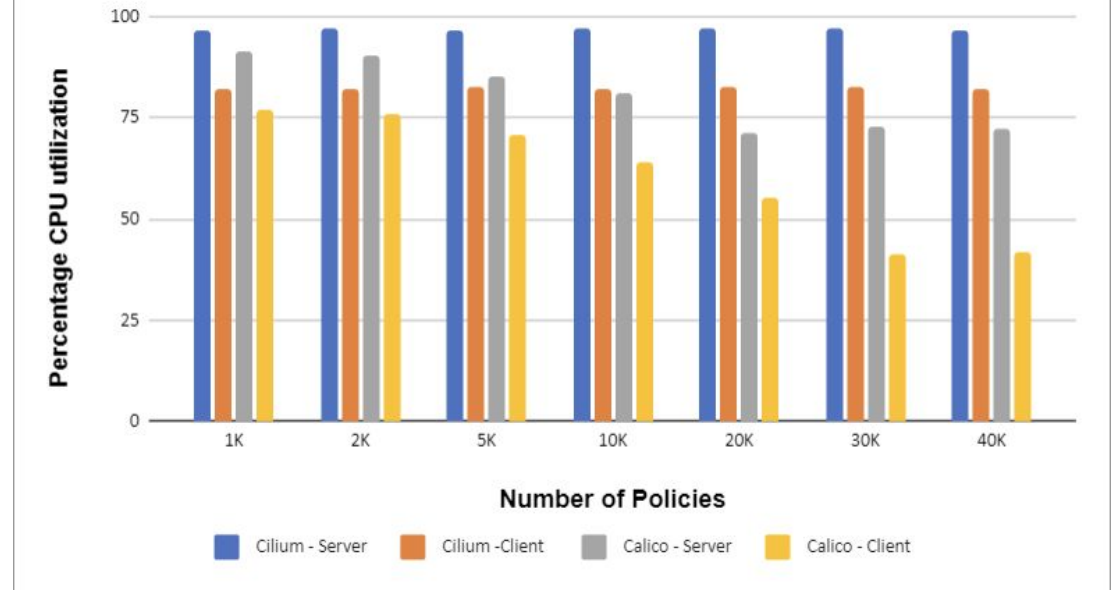
```

Layer3/4 policy testing for TCP

Layer 3 TCP Comparison - With increasing number of Policies



TCP - CPU Utilization as policies increase



Cilium eBPF Maps and Programs

```
an001@workercilium1:~$ sudo bpftool prog -p | grep cil*
```

```
"name": "cil_from_overlay",
"name": "cil_to_overlay",
"name": "cil_sock6_conne",
"name": "cil_sock6_post_",
"name": "cil_sock6_sendm",
"name": "cil_sock6_recvm",
"name": "cil_sock4_conne",
"name": "cil_sock4_post_",
"name": "cil_sock4_sendm",
"name": "cil_sock4_recvm",
"name": "cil_from_contai",
"name": "cil_to_host",
"name": "cil_from_host",
"name": "cil_to_host",
"name": "cil_from_netdev",
"name": "cil_to_netdev",
```

```
an001@workercilium1:~$ sudo bpftool map list -p | grep cil
```

```
"name": "cilium_lxc",
"name": "cilium_node_map",
"name": "cilium_metrics",
"name": "cilium_lb4_reve",
"name": "cilium_lb4_serv",
"name": "cilium_lb4_back",
"name": "cilium_lb4_reve",
"name": "cilium_events",
"name": "cilium_signals",
"name": "cilium_call_pol",
"name": "cilium_ct4_glob",
"name": "cilium_ct_any4_",
"name": "cilium_snat_v4_",
"name": "cilium_nodeport",
"name": "cilium_ipv4_fra",
"name": "cilium_lb_affin",
"name": "cilium_lb4_affi",
"name": "cilium_lb4_sour",
"name": "cilium_ipcache",
"name": "cilium_tunnel_m",
"name": "cilium_calls_ov",
"name": "cilium_encrypt_",
"name": "cilium_policy_0",
"name": "cilium_policy_0",
"name": "cilium_tail_cal",
"name": "cilium_calls_00",
"name": "cilium_calls_ho",
"name": "cilium_calls_ne",
"name": "cilium_calls_ne"
```

Calico IPTABLES entries

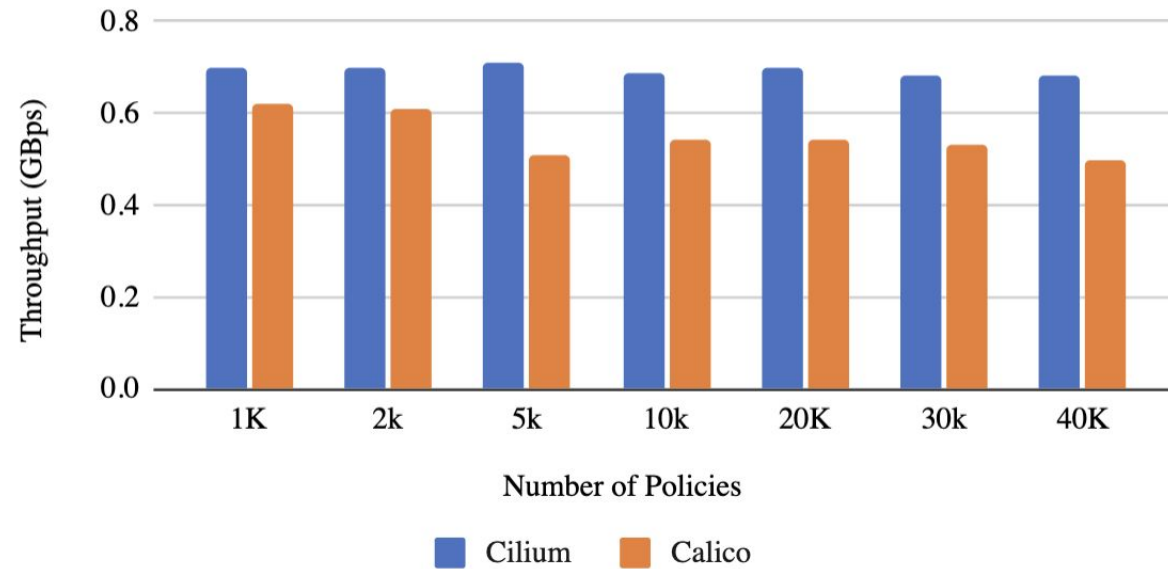
```
MARK      all -- anywhere anywhere /* cali:zR4GRLGUxP8e0JUT */ match-set cali40s:9BXRLmhJumopWvR8vk-q9Rl src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:cAitB3399eNLSdEP */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:u3FEexeyWiBMT6S4 */ match-set cali40s:378kgnyDD6BeIAL3eRhu_0 src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:cn8zUCjeCDurYCZg */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:raSdmXzY_1Q1UcE7 */ match-set cali40s:PxC0_TftDmCm4RNnjeY6Kld src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:lI2LpKhyvs7u60YM */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:psbY0nr8rL-LJpMG */ match-set cali40s:tiurfVgvw6PC9_fRMfFNDj6 src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:KwcUmYfQqIXdpS6A */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:x-jdw-cvkhR28FNW */ match-set cali40s:imjqvrEKMkpqLZrzKAu36FT src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:Hsb8pKM05IkpzdYw */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:cI2XfKEUeM9KQa9W */ match-set cali40s:2Yo6427DlrjvLGBseShyDRT src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:NBBQT0VAqnMKNgQV */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:Uq2NxvDH990r9JAag */ match-set cali40s:-9TvsX35emtms3xaCucNt_b src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:c3euPUlopL9_NJDj */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:0XA_73GwXz1iWeN */ match-set cali40s:nPC6-PVy3LaTgcaYlJqxedh src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:E2TwldID0-cNfjnT */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:fnVLBK88e48XjwQ0 */ match-set cali40s:R8990QM-Jmtqf2ymqoPgno src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:vU70kGRep7cfv9p */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:rRtMAwB7_v22L5AP */ match-set cali40s:NvdJqmqVs4DXxvCryBUAYnn src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:klBy4Bk0rkUyUahp */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:wJUz1mi3AgS-qydh */ match-set cali40s:uoojXlJUKwW-hgDb8lBP3lg src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:juqn2M1Aax5r-Mr0 */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:2F9t-PbdYnatWi8s */ match-set cali40s:kl2JV0wolaQPDeZu8EqNxtY src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:_MctbnUAdr0FwWE0 */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:pR5Jc4Yfy376vQCD */ match-set cali40s:ecnz1YKTdnT01FDvMP3Ht_W src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:t7yb_1s6gqGbvPyI */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:_BPYImLgBEGFaUI */ match-set cali40s:9lmszVRE28IGqAR_Pwb2QZe src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:XLCS09i2ZAtkHsIm */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:9RkA59qawRHY_TnV */ match-set cali40s:kseuzzoT8Eh-TX9nLiss9Jz src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:CrV5k0yo_fE8cd9d */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:mYq70XhbPqIITUGR */ match-set cali40s:67yAFdBu0DyWRu350kQ9Z89 src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:3PggrnVk-JtqchkL */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:Ig2_RAIdMwXLYuF */ match-set cali40s:yow-VicxTDMCPi9ja4a6sVU src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:UKP_7Dg7pu2JtITB */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:GpgtUW_4HphTd6cR */ match-set cali40s:dcqqBrVd4U9Plx6chQIwvcv src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:YQWphMyHxrf0fu21 */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:v01fUPB3ByNJ0ybh */ match-set cali40s:t1FqIDdRWZMcYpC_gtm4ton src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:GAHaE21A46nR_wb5 */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:FB0qnMLMSTLX45We */ match-set cali40s:KtN-ye9CP8cQ8pVvytHUs-B src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:zx7b4Y_yJpCsmvrg */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:zxgbTQaf0H8vKj_ */ match-set cali40s:edD8UT1lIo_XyzTj6K_qTww src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:DDgcRbf2dgsKPCf8 */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:Y-2HDIo_yxc05IuK */ match-set cali40s:5Pb5kkExMbKkuDf4WL-jRnw3 src MARK or 0x10000
RETURN    all -- anywhere anywhere /* cali:4B_etYv9nSSHkGwK */ mark match 0x10000/0x10000
MARK      all -- anywhere anywhere /* cali:KLr6nwFWghnq1md1 */ match-set cali40s:23tMYeP0ksLD5QceRY4aPD4 src MARK or 0x10000
```

Layer3/4 policy testing for TCP

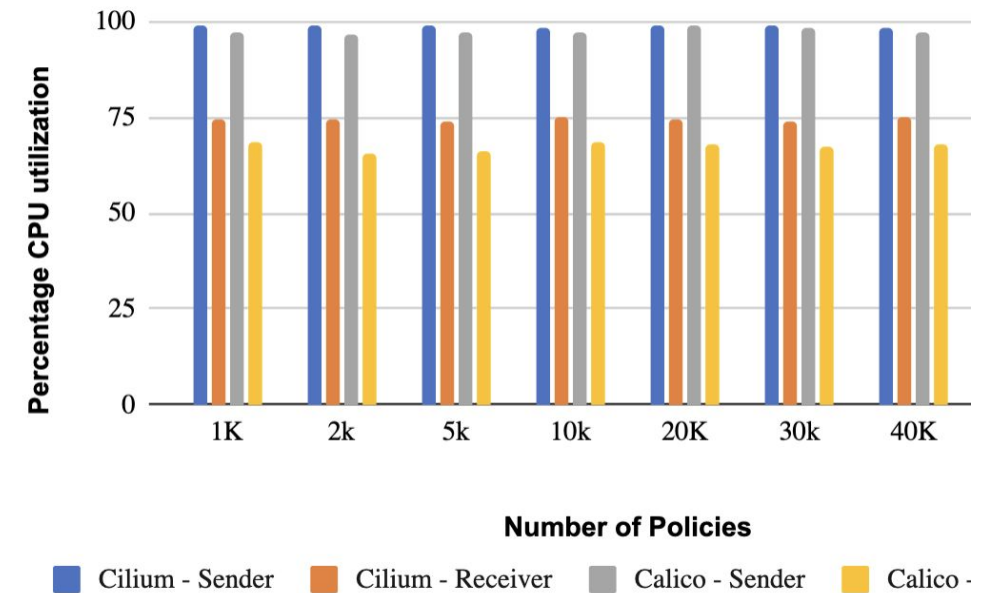
- Initial noticeable delay when starting the iperf experiment with the policies with Calico
- For Calico, the first few packets show least throughput but continue to improve slightly with time. This can be explained by the connection tracking Calico uses.
- Cilium throughput remains the same in both cases.

Layer3/4 based policy testing for UDP

Layer 3 UDP Comparison - With increasing number of Policies



UDP - CPU Utilization as policies increase



Scripts to Generate L7 Policies

```
CiliumProject > CiliumNetworkPolicy > $ CiliumpolicygeneratorLabel7.sh
1  #!/bin/bash
2  v="v17"
3  x=0
4  NUMBER=$1
5  rm ciliumbasepolicyv17.yaml
6  kubectl delete cnf iperfpolicyv300
7  kubectl delete cnf iperfpolicy$V
8
9  cat <<EOF > ciliumbasepolicyv17.yaml
10 kind: CiliumNetworkPolicy
11 apiVersion: "cilium.io/v2"
12 metadata:
13   name: "iperfpolicy$V"
14   namespace: default
15 spec:
16   endpointSelector:
17     matchLabels:
18       app: nginx
19       role: server
20   ingress:
21     - fromEndpoints:
22       - matchLabels:
23         app: apache
24         role: client
25       toPorts:
26         - ports:
27           - port: "80"
28           protocol: TCP
29         rules:
30           http:
31             - method: GET
32               path: "/get.html"
33             - method: POST
34               path: "/post.html"
35             - method: GET
36               path: "/"
37 EOF
38
39 while [ $x -lt $NUMBER ]; do
40   my_list=("GET" "PUT" "POST")
41   random_index=$(( $RANDOM % ${#my_list[@]} ))
42   length=5
43   y="${my_list[random_index]}"
44   random_string=$(head /dev/urandom | tr -dc 'a-zA-Z0-9' | head -c 5)
45   echo "    - method: $y" >> ciliumbasepolicyv17.yaml
46   echo "      path: \"/${random_string}.html\"" >> ciliumbasepolicyv17.yaml
47   ((x++))
48 done
49
50 echo "Uploading Policy !!!!!!!!!!!!!!!"
51 kubectl create -f ciliumbasepolicyv17.yaml
52
53 sleep 3
54 exit
```

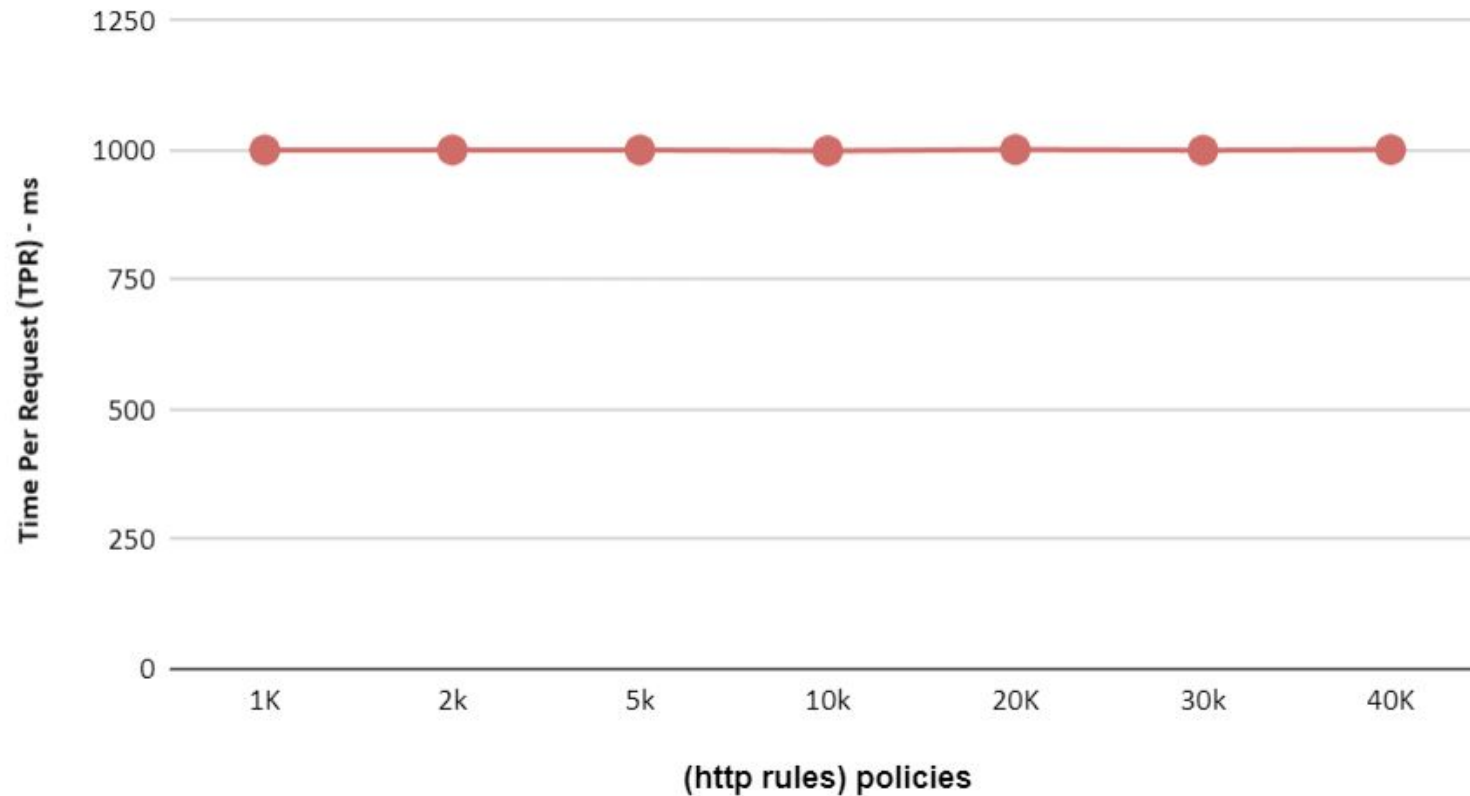
```
CiliumProject > CiliumNetworkPolicy > $ ciliumbasepolicyv17.yaml
1  kind: CiliumNetworkPolicy
2  apiVersion: "cilium.io/v2"
3  metadata:
4    name: "iperfpolicyv17"
5    namespace: default
6  spec:
7    endpointSelector:
8      matchLabels:
9        app: nginx
10       role: server
11    ingress:
12      - fromEndpoints:
13        - matchLabels:
14          app: apache
15          role: client
16        toPorts:
17          - ports:
18            - port: "80"
19            protocol: TCP
20          rules:
21            http:
22              - method: POST
23                path: "/8oRnc.html"
24              - method: POST
25                path: "/2FhMA.html"
26              - method: GET
27                path: "/XyRmi.html"
28              - method: POST
29                path: "/8BPed8.html"
30              - method: POST
31                path: "/s4ZEB.html"
32              - method: POST
33                path: "/uWd11.html"
34              - method: POST
35                path: "/20h5a.html"
36              - method: PUT
37                path: "/pT0at.html"
38              - method: GET
39                path: "/3ypXT.html"
40              - method: GET
41                path: "/f5ps8.html"
42              - method: PUT
43                path: "/hF119.html"
44              - method: POST
45                path: "/6SjAl.html"
46              - method: POST
47                path: "/eArXe.html"
48              - method: PUT
49                path: "/61Du.html"
50              - method: GET
51                path: "/v3fLO.html"
52              - method: GET
53                path: "/1Ww9.html"
54              - method: PUT
55                path: "/6F1ca.html"
56              - method: GET
57                path: "/EPW09.html"
58              - method: PUT
59                path: "/17mqX.html"
60              - method: PUT
61                path: "/ZT0L9.html"
62              - method: POST
63                path: "/napUy.html"
64              - method: POST
65                path: "/6Cn9.html"
66              - method: PUT
67                path: "/ki0ab.html"
68              - method: GET
69                path: "/7Eypn.html"
70              - method: PUT
71                path: "/bt0CN.html"
72              - method: GET
73                path: "/sARD5.html"
74              - method: PUT
75                path: "/ufz8d.html"
76              - method: GET
```

Apache-Benchmark Testing Script

```
1 #!/bin/bash
2 DATE=$(date +"%b_%d_%H%M%S")
3
4 kubectl apply -f abpod.yaml
5 sleep 3
6 kubectl apply -f nginxserver.yaml
7 sleep 3
8 kubectl cp get.html httpd-pod:/usr/local/apache2/htdocs/
9 kubectl cp post.html httpd-pod:/usr/local/apache2/htdocs/
10
11 PODNAME=$(kubectl get pods -o wide | grep apache-benchmark-pod | awk '{print $1}')
12 echo "Podname --> $PODNAME"
13
14 NGINX=$(kubectl get pods -o wide | grep httpd-pod | awk '{print $1}')
15 echo "iperfServer --> $NGINX "
16
17 HOSTIP=$(kubectl get pod $NGINX -o=jsonpath='{.status.podIP}')
18 echo "HostIP --> $HOSTIP"
19 echo "$DATE"
20
21 for i in {1..5}
22 do
23     echo "##### Simple request test $i #####" >> $DATE"Result".yaml
24     kubectl exec -it $PODNAME -- ab -n 1 -c 1 -i http://$HOSTIP:80/ | grep -E 'Transfer rate|Time per request|Requests per second' >> $DATE"Result".yaml
25     kubectl exec -it $PODNAME -- ab -n 1 -c 1 -i http://$HOSTIP:80/get.html | grep -E 'Transfer rate|Time per request|Requests per second' >> $DATE"Result".yaml
26     kubectl exec -it $PODNAME -- ab -n 1 -c 1 -i http://$HOSTIP:80/post.html | grep -E 'Transfer rate|Time per request|Requests per second' >> $DATE"Result".yaml
27     echo "#####" >> $DATE"Result".yaml
28     echo "#####" >> $DATE"Result".yaml
29     echo "#####" >> $DATE"Result".yaml
30 done
31
32
33 #CleanUP
34 echo "Clean UP in progress ....."
35
36 kubectl delete pods apache-benchmark-pod httpd-pod
37 sleep 2
38 echo "Test completes "
39 exit
```

Layer 7 - Policy Testing

Layer 7: Apache Benchmark



- a. Cilium uses hashing based L7 ip mapping
- b. consistent with our findings as TPR remains same for increasing number of requests

Summary of the Results

- Throughput: Cilium maintains a constant throughput as the number of policies increases, whereas Calico's throughput decreases due to iptables overhead.
- UDP Performance: Cilium demonstrates better performance in terms of packet loss and throughput consistency compared to Calico, which experiences higher packet loss.
- Time per request (without policies): Both Cilium and Calico have similar performance when there are no policies in place.
- L7 Policy Performance: Cilium is able to maintain constant throughput with an increasing number of L7 policies.

Conclusion

- Cilium's use of eBPF and XDP technologies provides fast and efficient packet filtering and forwarding, while its hash-based approach to L7 IP rule mapping allows for efficient rule matching.
- These results suggest that Cilium's eBPF-based approach provides significant performance advantages over traditional CNIs like Calico, particularly in terms of throughput consistency, UDP performance, and L7 policy handling.

Future work

- Perform similar testing with increased number of cores. This can help improve performance - cpu utilization, reduce packet loss, etc
- Perform testing directly on instances (bare metal) compared to VMs and compare results of both
- Use custom built linux kernels which allows to more finely tune network configurations

References

- [eBPF](#)
- [CNI Benchmark: Understanding Cilium Network Performance](#)
- Assessing Container Network Interface Plugins: Functionality, Performance, and Scalability -
<https://ieeexplore.ieee.org/document/9309003>
- [Performance Benchmarking and Tuning for Container Networking on Arm](#)
- Cilium Documentation : [Performance Evaluation – Cilium 1.9.15 documentation](#)
- Similar experimentation:
<https://kinvolk.io/blog/2020/12/egress-filtering-benchmark-part-2-calico-and-cilium/>
- Github : <https://github.com/anvayabn/CiliumProject>
-

Thank you